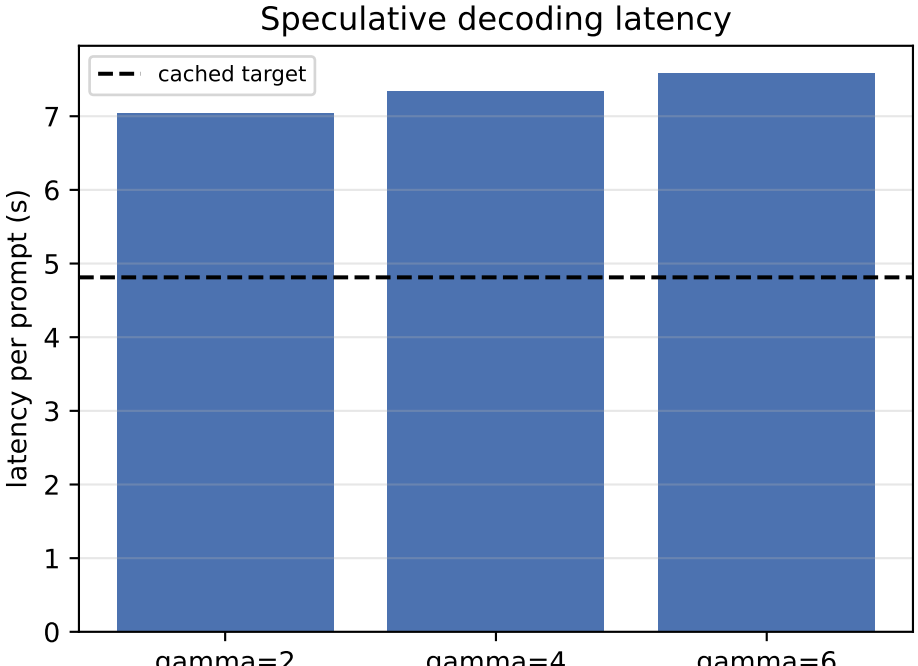
Speculative decoding latency
cached target
latency per prompt (s)
0
1
2
3
4
5
6
7
gamma=2
gamma=4
gamma=6

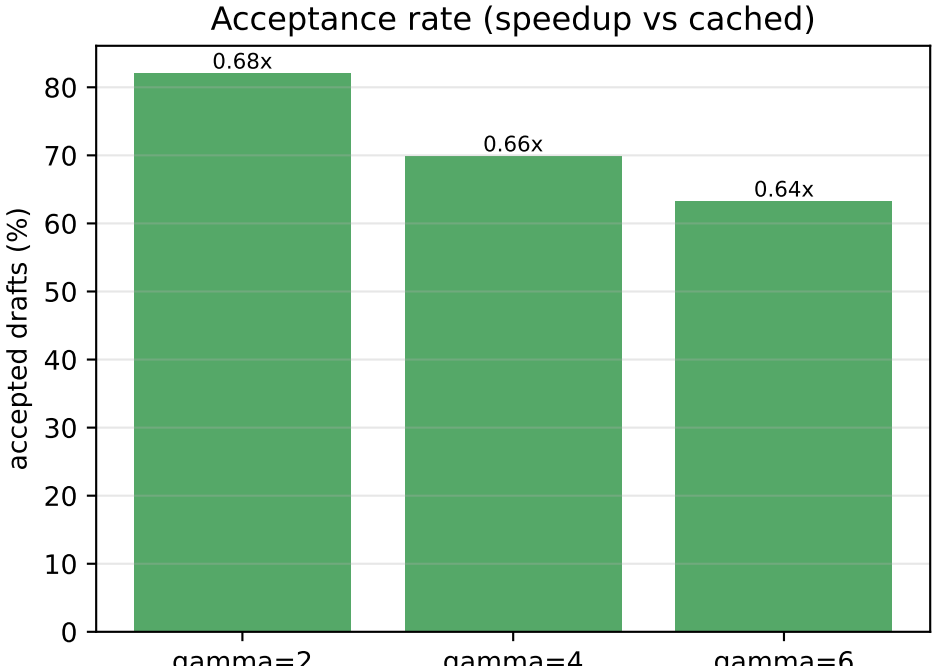
Acceptance rate (speedup vs cached)
accepted drafts (%)
0
10
20
30
40
50
60
70
80
0.68x
0.66x
0.64x
gamma=2
gamma=4
gamma=6